


How good is the invasion model? (h = 1w)

(A) calibration + AUC and (B) real-time prioritisation vs random – after Kraemer & Cauchemez 2017, Lancet Inf Dis, Fig 3

(A) Calibration: predicted vs observed invasion (h = 1w)


On the diagonal = calibrated probabilities; below = over-confident.



The 3.2% of at-risk zone-weeks we score >15%/week had a 18% observed invasion rate.

(B) Prioritisation: invasions caught vs random targeting

Share of true next-week invasions caught at a fixed weekly alert budget of K zones, pooled over LFO folds



At a budget of K=10 zones/week, Bayes-M10-med catches 55% of the next invasions vs 2% for random targeting (26.9x better).