


Of the top-K highest-risk zones, what % were actually invaded? (h=1w) – Bayesian models

Precision at a fixed weekly alert budget of K zones, pooled over LFO folds; dashed = random-watch-list base rate



Higher = fewer false alarms. Precision declines with K as the budget reaches lower-risk zones; the gap above the base rate is the model's targeting value